



Raw reward model scores during RL

lines: 25-rollout trailing mean · no length penalty in this scalar · compare within each RM family


Qwen3-14B Base → SFT → RL



Qwen3-14B Instruct → RL



Llama-3.1-8B Base → SFT → RL



Llama-3.1-8B Instruct → RL

