



Response length under RL — {Qwen, Llama} × {base, instruct}

lines: 25-rollout trailing mean · soft length reward active in all four suites (8/1024/2048 window, shared σ_0 per suite)


Qwen3-14B Base → SFT → RL



Qwen3-14B Instruct → RL



Llama-3.1-8B Base → SFT → RL



Llama-3.1-8B Instruct → RL

